


Analytics 3.0

by Thomas H. Davenport

FROM THE DECEMBER 2013 ISSUE



ARTWORK: CHAD HAGEN, NONSENSICAL INFOGRAPHIC NO. 5, 2009, DIGITAL

Those of us who have spent years studying “data smart” companies believe we’ve already lived through two eras in the use of analytics. We might call them BBD and ABD—before big data and after big data. Or, to use a naming convention matched to the topic, we might say that Analytics 1.0 was followed by Analytics 2.0. Generally speaking, 2.0 releases don’t just add some bells and whistles or make minor performance tweaks. In contrast to, say, a 1.1 version, a 2.0 product is a more substantial overhaul based on new priorities and technical possibilities. When large numbers of companies began capitalizing on vast new sources of unstructured, fast-moving information—big data—that was surely the case.

Some of us now perceive another shift, fundamental and far-reaching enough that we can fairly call it Analytics 3.0. Briefly, it is a new resolve to apply powerful data-gathering and analysis methods not just to a company’s operations but also to its offerings—to embed data smartness into the products and services customers buy.

I’ll develop this argument in what follows, making the case that just as the early applications of big data marked a major break from the 1.0 past, the current innovations of a few industry leaders are evidence that a new era is dawning. When a new way of thinking about and applying a strength begins to take hold, managers are challenged to respond in many ways. Change comes fast to every part of a business’s world. New players emerge, competitive positions shift, novel technologies must be mastered, and talent gravitates toward the most exciting new work.

Managers will see all these things in the coming months and years. The ones who respond most effectively will be those who have connected the dots and recognized that competing on analytics is being rethought on a large scale. Indeed, the first companies to perceive the general direction of change—those with a sneak peek at Analytics 3.0—will be best positioned to drive that change.

The Evolution of Analytics

My purpose here is not to make abstract observations about the unfolding history of analytics. Still, it is useful to look back at the last big shift and the context in which it occurred. The use of data to make decisions is, of course, not a new idea; it is as old as decision making itself. But the field of business analytics was born in the mid-1950s, with the advent of tools that could produce and capture a larger quantity of information and discern patterns in it far more quickly than the unassisted human mind ever could.

Today it isn’t just online and information firms that can create products and services from analyses of data. It’s every firm in every industry.

Analytics 1.0—the era of “business intelligence.”

What we are here calling Analytics 1.0 was a time of real progress in gaining an objective, deep understanding of important business phenomena and giving managers the fact-based comprehension to go beyond intuition when making decisions. For the first time, data about production processes, sales, customer interactions, and more were recorded, aggregated, and analyzed.

New computing technologies were key. Information systems were at first custom-built by companies whose large scale justified the investment; later, they were commercialized by outside vendors in more-generic forms. This was the era of the enterprise data warehouse, used to capture information, and of business intelligence software, used to query and report it.

New competencies were required as well, beginning with the ability to manage data. Data sets were small enough in volume and static enough in velocity to be segregated in warehouses for analysis. However, readying a data set for inclusion in a warehouse was difficult. Analysts spent much of their time *preparing* data for analysis and relatively little time on the analysis itself.

More than anything else, it was vital to figure out the right few questions on which to focus, because analysis was painstaking and slow, often taking weeks or months to perform. And reporting processes—the great majority of business intelligence activity—addressed only what had happened in the past; they offered no explanations or predictions.

Did people see analytics as a source of competitive advantage? In broad terms, yes—but no one spoke in today’s terms of “competing on analytics.” The edge came in the form of greater operational efficiency—making better decisions on certain key points to improve performance.

Analytics 2.0—the era of big data.

The basic conditions of the Analytics 1.0 period predominated for half a century, until the mid-2000s, when internet-based and social network firms primarily in Silicon Valley—Google, eBay, and so on—began to amass and analyze new kinds of information. Although the term “big data” wasn’t coined immediately, the new reality it signified very quickly changed the role of data and analytics in those firms. Big data also came to be distinguished from small data because it was not generated purely by a firm’s internal transaction systems. It was externally sourced as well, coming from the internet, sensors of various types, public data initiatives such as the human genome project, and captures of audio and video recordings.

As analytics entered the 2.0 phase, the need for powerful new tools—and the opportunity to profit by providing them—quickly became apparent. Companies rushed to build new capabilities and acquire customers. The broad recognition of the advantage a first mover could gain led to an impressive level of hype but also prompted an unprecedented acceleration of new offerings. LinkedIn, for example, has created numerous data products, including People You May Know, Jobs You May Be Interested In, Groups You May Like, Companies You May Want to Follow, Network Updates, and Skills and Expertise. To do so, it built a strong infrastructure and hired smart, productive data scientists. Its highly successful Year in Review, which summarizes the job changes of people in a member’s network, was developed in just a month. And LinkedIn is not the only company focused on speed. One CEO of a big data start-up told me, “We tried agile [development methodology], but it was too slow.”

Innovative technologies of many kinds had to be created, acquired, and mastered. Big data couldn’t fit or be analyzed fast enough on a single server, so it was processed with Hadoop, an open source software framework for fast batch data processing across parallel servers. To deal with relatively unstructured data, companies turned to a new class of databases known as NoSQL. Much information was stored and analyzed in public or private cloud-computing environments. Other technologies introduced during this period include “in memory” and “in database” analytics for fast number crunching. Machine-learning methods (semiautomated model development and testing) were used to rapidly generate models from the fast-moving data. Black-and-white reports gave way to colorful, complex visuals.

Thus, the competencies required for Analytics 2.0 were quite different from those needed for 1.0. The next-generation quantitative analysts were called data scientists, and they possessed both computational and analytical skills. Soon the data scientists were not content to remain in the back office; they wanted to work on new product offerings and help shape the business.

Analytics 3.0—the era of data-enriched offerings.

During 2.0, a sharp-eyed observer could have seen the beginnings of analytics’ next big era. The pioneering big data firms in Silicon Valley began investing in analytics to support customer-facing products, services, and features. They attracted viewers to their websites through better search algorithms, recommendations from friends and colleagues, suggestions for products to buy, and highly targeted ads, all driven by analytics rooted in enormous amounts of data.

FURTHER READING

Analytics 3.0 marks the point when other large organizations started to follow suit. Today it’s not just information firms and online companies that can create products and services from analyses of data. It’s every firm in



Making Advanced Analytics Work for You

ANALYTICS FEATURE by Dominic Barton and David Court

A guide for managers who want to get started with big data.

SAVE SHARE

every industry. If your company makes things, moves things, consumes things, or works with customers, you have increasing amounts of data on those activities. Every device, shipment, and consumer leaves a trail. You have the ability to analyze those sets of data for the benefit of customers and markets. You also have the ability to embed analytics and optimization into every business decision made at the front lines of your operations.

Like the first two eras of analytics, this one brings new challenges and opportunities, both for the companies that want to compete on analytics and for the vendors that supply the data and tools with which to do so. How to capitalize on the shift is a subject we will turn to shortly. First, however, let's consider what Analytics 3.0 looks like in some well-known firms—all of which were decidedly offline businesses for most of their many decades in operation.

The Next Big Thing, in Beta

The Bosch Group, based in Germany, is 127 years old, but it's hardly last-century in its application of analytics. The company has embarked on a series of initiatives across business units that make use of data and analytics to provide so-called intelligent customer offerings. These include intelligent fleet management, intelligent vehicle-charging infrastructures, intelligent energy management, intelligent security video analysis, and many more. To identify and develop these innovative services, Bosch created a Software Innovations group that focuses heavily on big data, analytics, and the "Internet of Things."

Schneider Electric, a 170-year-old company based in France, originally manufactured iron, steel, and armaments. Today it focuses primarily on energy management, including energy optimization, smart-grid management, and building automation. It has acquired or developed a variety of software and data ventures in Silicon Valley, Boston, and France. Its Advanced Distribution Management System, for example, handles energy distribution in utility companies. ADMS monitors and controls network devices, manages service outages, and dispatches crews. It gives utilities the ability to integrate millions of data points on network performance and lets engineers use visual analytics to understand the state of the network.

One of the most dramatic conversions to data and analytics offerings is taking place at General Electric, a company that's more than 120 years old. GE's manufacturing businesses are increasingly becoming providers of asset and operations optimization services. With sensors streaming data from turbines, locomotives, jet engines, and medical-imaging devices, GE can determine the most efficient and effective service intervals for those machines. To assemble and develop the skilled employees needed for this work, the company invested more than \$2 billion in a new software and analytics center in the San Francisco Bay area. It is now selling technology to other industrial companies for use in managing big data and analytics, and it has created new technology offerings based on big data concepts, including Predix (a platform for building "industrial internet" applications) and Predictivity (a series of 24 asset or operations optimization applications that run on the Predix platform across industries).

UPS, a mere 107 years old, is perhaps the best example of an organization that has pushed analytics out to frontline processes—in its case, to delivery routing. The company is no stranger to big data, having begun tracking package movements and transactions in the 1980s. It captures information on the 16.3 million packages, on average, that it delivers daily, and it receives 39.5 million tracking requests a day. The most recent source of big data at UPS is the telematics sensors in more than 46,000 company trucks, which track metrics including speed, direction, braking, and drivetrain performance. The waves of incoming data not only show daily performance but also are informing a major redesign of drivers' routes. That initiative, called ORION (On-Road Integrated Optimization and Navigation), is arguably the world's largest operations research project. It relies heavily on online map data and optimization algorithms and will eventually be able to reconfigure a driver's pickups and deliveries in real time. In 2011 it cut 85 million miles out of drivers' routes, thereby saving more than 8.4 million gallons of fuel.

The common thread in these examples is the resolve by a company's management to compete on analytics not only in the traditional sense (by improving internal business decisions) but also by creating more-valuable products and services. This is the essence of Analytics 3.0.

Some readers will recognize the coming era as the realization of a prediction made long ago. In their 1991 book *2020 Vision*, Stan Davis and Bill Davidson argued that companies should “informationalize” their businesses—that is, develop products and services on the basis of information. They observed that companies emit “information exhaust” that could be captured and used to “turbocharge” their offerings. At the time, their ideas gained traction only among companies already in the information business, such as Quotron (stock data) and the Official Airline Guide (flight data). But today banks, industrial manufacturers, health care providers, retailers—any company, in any industry, that is willing to exploit the possibilities—can develop valuable products and services from their aggregated data.

Google, Amazon, and others have prospered not by giving customers information but by giving them shortcuts to decisions and actions.

Davis and Davidson wrote at a time when supplying information was enough. But these days we are inundated with information and have little time to turn it into insight. Companies that were information providers must become insight providers, using analytics to digest information and tell us what to do with it. Online businesses, with vast amounts of clickstream data at their disposal, have pioneered this approach: Google, LinkedIn, Facebook, Amazon, and others have prospered not by giving customers information but by giving them shortcuts to decisions and actions. Companies in the conventional information industry are now well along this path too.

Ten Requirements for Capitalizing on Analytics 3.0

This strategic change in focus means a new role for analytics within organizations. Companies will need to recognize a host of related challenges and respond with new capabilities, positions, and priorities.

Multiple types of data, often combined.



Organizations will need to integrate large and small volumes of data from internal and external sources and in structured and unstructured formats to yield new insights in predictive and prescriptive models—ones that tell frontline workers how best to perform their jobs. The trucking company Schneider National, for example, is adding data from new sensors to its logistical optimization algorithms, allowing it to monitor key indicators such as fuel levels, container location and capacity, and driver behavior. It aims to steadily improve the efficiency of its route networks, lower its fuel costs, and decrease the risk of accidents.

BY THE SAME AUTHOR

Competing on Analytics

ANALYTICS FEATURE by Thomas H. Davenport

How to compete on the quantitative turf.

 SAVE  SHARE

A new set of data management options.

In the 1.0 era, firms used data warehouses as the basis for analysis. In the 2.0 era, they focused on Hadoop clusters and NoSQL databases. Today the technology answer is “all of the above”: data warehouses, database and big data appliances, environments that combine traditional data query approaches with Hadoop (these are sometimes called Hadoop 2.0), vertical

and graph databases, and more. The number and complexity of choices IT architects must make about data management have expanded considerably, and almost every organization will end up with a hybrid data environment. The old formats haven’t gone away, but new processes are needed to move data and analysis across staging, evaluation, exploration, and production applications.

Faster technologies and methods of analysis.

Big data technologies from the 2.0 period are considerably faster than previous generations of technology for data management and analysis were. To complement them, new “agile” analytical methods and machine-learning techniques are being used to produce insights at a much faster rate. Like agile systems development, these methods involve frequent delivery of partial outputs to the project stakeholders; as with the best data scientists’ work, they have an ongoing sense of urgency. The challenge in the 3.0 era is to adapt operational, product development, and decision processes to take advantage of what the new technologies and methods can bring forth.

Embedded analytics.

Consistent with the increased speed of data processing and analysis, models in Analytics 3.0 are often embedded into operational and decision processes, dramatically increasing their speed and impact. For example, Procter & Gamble is integrating analytics in day-to-day management decision making through more than 50 “business sphere” decision rooms and more than 50,000 “decision cockpits” on employee computers.

Some firms are embedding analytics into fully automated systems through scoring algorithms and analytics-based rules. Some are building analytics into consumer-oriented products and features. Whatever the scenario, integrating analytics into systems and processes not only means greater speed but also makes it harder for decision makers to avoid using analytics—which is usually a good thing.

Imagine if Ron Johnson’s tenure at J.C. Penney had involved small-scale, data-driven experiments rather than wholesale changes.

Data discovery.

To develop products and services on the basis of data, companies need a capable discovery platform for data exploration along with the requisite skills and processes. Although enterprise data warehouses were initially intended to facilitate exploration and analysis, they have become production data repositories for many organizations, and, as previously noted, getting data into them is time-consuming. Data discovery environments make it possible to determine the essential features of a data set without a lot of preparation.

Cross-disciplinary data teams.

In online firms and big data start-ups, data scientists are often able to run the whole show (or at least to have a lot of independence). In larger and more conventional firms, however, they must collaborate with a variety of other players to ensure that big data is matched by big analytics. In many cases the “data scientists” in such firms are actually conventional quantitative analysts who are forced to spend a bit more time than they’d like on data management activities (hardly a new phenomenon). Companies now employ data hackers, who excel at extracting and structuring information, to work with analysts, who excel at modeling it.

Both groups have to work with IT, which supplies the big data and the analytical infrastructure, provisions the “sandboxes” in which the groups explore the data, and turns exploratory analysis into production capabilities. The combined team takes on whatever is needed to get the analytical job done, with frequent overlap among roles.

Chief analytics officers.

When analytics are this important, they need senior management oversight. Companies are beginning to create “chief analytics officer” roles to superintend the building and use of analytical capabilities. Organizations with C-level analytics leaders include AIG, FICO, USAA, the University of Pittsburgh Medical Center, the Obama reelection campaign, Wells Fargo, and Bank of America. The list will undoubtedly grow.

Prescriptive analytics.

There have always been three types of analytics: descriptive, which reports on the past; predictive, which uses models based on past data to predict the future; and prescriptive, which uses models to specify optimal behaviors and actions. Although Analytics 3.0 includes all three types, it emphasizes the last. Prescriptive models involve large-scale testing and optimization and are a means of embedding analytics into key processes and employee behaviors. They provide a high level of operational benefits but require high-quality planning and execution in return. For example, if the UPS ORION system gives incorrect routing information to drivers, it won’t be around for long. UPS executives say they have spent much more time on change management issues than on algorithm and systems development.

Analytics on an industrial scale.

For companies that use analytics mainly for internal decision processes, Analytics 3.0 provides an opportunity to scale those processes to industrial strength. Creating many more models through machine learning can let an organization become much more granular and precise in its predictions. IBM, for instance, formerly used 150 models in its annual “demand generation” process, which assesses which customer accounts are worth greater investments of salesperson time and energy. Working with a small company, Modern Analytics, and using a “model factory” and “data assembly line” approach, IBM now creates and maintains 5,000 such models a year—and needs just four people to do so. Its new systems can build 95% of its models without any human intervention, and another 3% require only minimal tuning from an analyst. And the new models address highly specific products, customer segments, and geographies. A test conducted in one large Asian market showed that such models doubled customer response rates compared with nonstatistical segmentation approaches.

New ways of deciding and managing.

In order for analytics to power the data economy in your company, you’ll need new approaches to decision making and management. Many will give you greater certainty before taking action. Managers need to become comfortable with data-driven experimentation. They should demand that any important initiative be preceded by small-scale but systematic experimentation of this sort, with rigorous controls to permit the determination of cause and effect. Imagine, for example, if Ron Johnson’s tenure as CEO of J.C. Penney had involved limited experiments rather than wholesale changes, most of which turned out badly.

Paradoxically, some of the changes prompted by the widespread availability of big data will not yield much certainty. Big data flows continuously—consider the analysis of brand sentiment derived from social media sources—and so metrics will inevitably rise and fall over time. Such “digital smoke signals,” as they have been called, can serve as an early warning system for budding problems. But they are indicative, not confirmatory. Managers will have to establish guidelines for when early warnings should cue decisions and action.

Additional uncertainty arises from the nature of big data relationships. Unless they are derived from formal testing, the results from big data generally involve correlation, not causation, and sometimes they occur by chance (although having greater amounts of data increases the likelihood that weak results will be statistically significant). Some managers may be frustrated by these facts. If the issue under consideration is highly important, further investigation may be warranted before a decision is made.

The use of prescriptive analytics often requires changes in the way frontline workers are managed. Companies will gain unprecedented visibility into the activities of truck drivers, airline pilots, warehouse workers, and any other employees wearing or carrying sensors (perhaps this means all employees, if smartphone sensors are included). Workers will undoubtedly be sensitive to this monitoring. Just as analytics that are intensely revealing of customer behavior have a certain “creepiness” factor, overly detailed reports of employee activity can cause discomfort. In the world of Analytics 3.0, there are times we need to look away.

Creating Value in the Data Economy

Does Analytics 3.0 represent the ultimate form of competing on analytics? Perhaps not. But it seems safe to say that it will be viewed as the point in time when participation in the data economy went mainstream.


The online companies that unleashed big data on the world were built around it from the beginning. They didn’t need to reconcile or integrate big data with traditional sources of information and the analytics performed on it, because for the most part, they didn’t have those traditional sources. They didn’t need to merge big data technologies with traditional IT infrastructures; in their companies, those infrastructures didn’t exist. Big data could stand alone, big data analytics could be the only analytics, and big data technology architectures could be the only IT architectures. But each of these companies now has its own version of Analytics 3.0.

One thing is clear: The new capabilities required of both long-established and start-up firms can’t be developed using old models for how analytics supported the business. The big data model was a huge step forward, but it will not provide advantage for much longer. Companies that want to prosper in the new data economy must once again fundamentally rethink how the analysis of data can create value for themselves and their customers. Analytics 3.0 is the direction of change and the new model for competing on analytics.

A version of this article appeared in the December 2013 issue of *Harvard Business Review*.

Thomas H. Davenport is a distinguished professor at Babson College, a research fellow at the MIT Center for Digital Business, and a senior adviser to Deloitte Analytics. He is at work on a book about automation in knowledge work.

This article is about ANALYTICS

 FOLLOW THIS TOPIC

Related Topics: DATA | STRATEGY | CUSTOMERS

Comments

Leave a Comment

POST

0 COMMENTS

✓ JOIN THE CONVERSATION

POSTING GUIDELINES

We hope the conversations that take place on HBR.org will be energetic, constructive, and thought-provoking. To comment, readers must sign in or register. And to ensure the quality of the discussion, our moderating team will review all comments and may edit them for clarity, length, and relevance. Comments that are overly promotional, mean-spirited, or off-topic may be deleted per the moderators' judgment. All postings become the property of Harvard Business Publishing.